

Exploration_ProspersLoan

October 7, 2020

1 Prosper Loan Data Analysis & Visualization

1.1 by Felicia Demulling

1.2 Preliminary Wrangling

The Prosper loan data set contains 113,937 loans with 81 variables on each loan, including loan amount, borrower rate (or interest rate), current loan status, borrower income, and many others. The loan information is spread throughout the United States.

```
In [251]: # import all packages and set plots to be embedded inline
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb

%matplotlib inline
```

```
In [252]: df = pd.read_csv('prosperLoanData.csv')
```

1.2.1 What is the structure of your dataset?

There are 81 columns in this dataset. Here are the columns that I will be using the most.

- ListingKey - Unique key for each listing
- ListingCreationDate - The date the listing was created
- Term - Length of the loan expressed in months
- LoanStatus - Status of the loan - canceled, current, defaulted, etc.
- BorrowerRate - Borrower's annual interest percentage rate
- Listing Category - what type of loan it is
- BorrowerState - the state where the listing was created
- Occupation: The occupation selected by the borrower
- CreditScoreRangeLower - The lower credit score
- CreditScoreRangeHigher - The upper credit score
- DebtToIncomeRatio - Debt to income ratio of borrower
- IncomeRange - Income of the borrower
- LoanOriginalAmount - The amount of the loan

1.2.2 What is/are the main feature(s) of interest in your dataset?

Being this dataset is so large, I have endless opportunities to find relationships. Here are the top most interesting things I'd like to explore.

- The average interest rate of most borrowers.
- What the highest, lowest and average risk scores are.
- What the most common loan types are.
- Which states have the highest borrowers.
- What are the different occupations and how much did they borrow?

I'm interested in the individuals income range compared to how much they borrowed.

1.2.3 What features in the dataset do you think will help support your investigation into your feature(s) of interest?

The features that I'll be using in the dataset the most are: the borrower rate will help me investigate how much each individual was charged for interest rate; the prosper score will allow me to see the average risk scores (based on a scale of 1 to 10, 10 being the

2 Clean

2.0.1 Step 1. Delete all unused columns

```
In [253]: df.drop(columns=['ListingNumber', 'ListingNumber', 'BorrowerAPR', 'LenderYield',  
                        'EstimatedEffectiveYield', 'EstimatedLoss', 'EstimatedReturn'], axis=1, inplace=True)  
  
df.drop(columns=['ProsperRating (numeric)', 'ProsperRating (Alpha)', 'EmploymentStatus',  
                'EmploymentStatusDuration', 'IsBorrowerHomeowner', 'CurrentlyInGroup'], axis=1, inplace=True)  
  
df.drop(columns=['GroupKey', 'DateCreditPulled', 'FirstRecordedCreditLine', 'CurrentCreditLines',  
                'OpenCreditLines', 'TotalCreditLinespast7years', 'OpenRevolvingAccounts'], axis=1, inplace=True)  
  
df.drop(columns=['OpenRevolvingMonthlyPayment', 'InquiriesLast6Months', 'TotalInquiriesLast6Months',  
                'CurrentDelinquencies', 'AmountDelinquent', 'DelinquenciesLast7Years'], axis=1, inplace=True)  
  
df.drop(columns=['PublicRecordsLast10Years', 'PublicRecordsLast12Months', 'RevolvingCreditUtilization',  
                'BankcardUtilization', 'AvailableBankcardCredit', 'TotalTrades'], axis=1, inplace=True)  
  
df.drop(columns=['TradesOpenedLast6Months', 'IncomeVerifiable', 'LoanKey', 'TotalProsperPaymentsBilled',  
                'TotalProsperPaymentsBilled', 'OnTimeProsperPayments', 'ProsperPaymentsLessThan30DaysLate'], axis=1, inplace=True)  
  
df.drop(columns=['ProsperPaymentsOneMonthPlusLate', 'ProsperPrincipalBorrowed', 'ProsperScoreChangeAtTimeOfListing',  
                'ScoreChangeAtTimeOfListing', 'LoanCurrentDaysDelinquent', 'LoanFirstDefaulted'], axis=1, inplace=True)  
  
df.drop(columns=['LoanMonthsSinceOrigination', 'LoanOriginationDate', 'LoanOriginationAmount', 'MonthlyLoanPayment',  
                'LP_CustomerPayments', 'LP_CustomerPrincipalPayments', 'LP_CustomerInterestPayments'], axis=1, inplace=True)
```

```
df.drop(columns=['LP_ServiceFees', 'LP_CollectionFees', 'LP_GrossPrincipalLoss', 'LP_NonPrincipalRecoverypayments', 'PercentFunded', 'Recommendations', 'InvestmentFromFriendsAmount', 'Investors', 'LoanNumber', 'TradesNever'])

df.head()
```

```
Out[253]:
```

	ListingKey	ListingCreationDate	CreditGrade	Term	\
0	1021339766868145413AB3B	2007-08-26 19:09:29.263000000	C	36	
1	10273602499503308B223C1	2014-02-27 08:28:07.900000000	NaN	36	
2	0EE9337825851032864889A	2007-01-05 15:00:47.090000000	HR	36	
3	0EF5356002482715299901A	2012-10-22 11:02:35.010000000	NaN	36	
4	0F023589499656230C5E3E2	2013-09-14 18:38:39.097000000	NaN	36	

	LoanStatus	ClosedDate	BorrowerRate	ProsperScore	\
0	Completed	2009-08-14 00:00:00	0.1580	NaN	
1	Current	NaN	0.0920	7.0	
2	Completed	2009-12-17 00:00:00	0.2750	NaN	
3	Current	NaN	0.0974	9.0	
4	Current	NaN	0.2085	4.0	

	ListingCategory (numeric)	BorrowerState	Occupation	\
0	0	CO	Other	
1	2	CO	Professional	
2	0	GA	Other	
3	16	GA	Skilled Labor	
4	2	MN	Executive	

	CreditScoreRangeLower	CreditScoreRangeUpper	DebtToIncomeRatio	\
0	640.0	659.0	0.17	
1	680.0	699.0	0.18	
2	480.0	499.0	0.06	
3	800.0	819.0	0.15	
4	680.0	699.0	0.26	

	IncomeRange	StatedMonthlyIncome	LoanOriginalAmount
0	\$25,000-49,999	3083.333333	9425
1	\$50,000-74,999	6125.000000	10000
2	Not displayed	2083.333333	3001
3	\$25,000-49,999	2875.000000	10000
4	\$100,000+	9583.333333	15000

```
In [254]: duplicate_df = df[df.duplicated(subset='ListingKey')]
```

Test

```
In [255]: duplicate_df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 871 entries, 9 to 113863
Data columns (total 17 columns):
ListingKey                871 non-null object
ListingCreationDate        871 non-null object
CreditGrade               0 non-null object
Term                      871 non-null int64
LoanStatus                871 non-null object
ClosedDate                13 non-null object
BorrowerRate              871 non-null float64
ProsperScore              871 non-null float64
ListingCategory (numeric) 871 non-null int64
BorrowerState             871 non-null object
Occupation                812 non-null object
CreditScoreRangeLower     871 non-null float64
CreditScoreRangeUpper    871 non-null float64
DebtToIncomeRatio         789 non-null float64
IncomeRange               871 non-null object
StatedMonthlyIncome       871 non-null float64
LoanOriginalAmount        871 non-null int64
dtypes: float64(6), int64(3), object(8)
memory usage: 122.5+ KB

```

```

In [256]: df.drop_duplicates(subset=['ListingKey'], keep='last', inplace=True)
          df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 113066 entries, 0 to 113936
Data columns (total 17 columns):
ListingKey                113066 non-null object
ListingCreationDate        113066 non-null object
CreditGrade               28953 non-null object
Term                      113066 non-null int64
LoanStatus                113066 non-null object
ClosedDate                55076 non-null object
BorrowerRate              113066 non-null float64
ProsperScore              83982 non-null float64
ListingCategory (numeric) 113066 non-null int64
BorrowerState             107551 non-null object
Occupation                109537 non-null object
CreditScoreRangeLower     112475 non-null float64
CreditScoreRangeUpper    112475 non-null float64
DebtToIncomeRatio         104594 non-null float64
IncomeRange               113066 non-null object
StatedMonthlyIncome       113066 non-null float64
LoanOriginalAmount        113066 non-null int64
dtypes: float64(6), int64(3), object(8)

```

memory usage: 15.5+ MB

2.0.2 Step 2. Replace category numbers with the name

```
In [257]: df.replace({'ListingCategory (numeric)': {0: 'Not Available', 1: 'Debt Consolidation',
3: 'Business', 4: 'Personal Loan', 5: 'Student Loan', 6: 'Car Loan', 7: 'Other', 8: 'Baby&Adoption', 9: 'Boat', 10: 'Mortgage',
11: 'Engagement Ring', 12: 'Green Loans', 13: 'Home Improvement', 14: 'Large Purchases', 15: 'Medical/Dental',
16: 'Personal Loan', 17: 'RV', 18: 'Taxes', 19: 'Vacation', 20: 'Wedding'}}
```

Test

```
In [258]: df.head()
```

```
Out[258]:
```

	ListingKey	ListingCreationDate	CreditGrade	Term	\
0	1021339766868145413AB3B	2007-08-26 19:09:29.263000000	C	36	
1	10273602499503308B223C1	2014-02-27 08:28:07.900000000	NaN	36	
2	0EE9337825851032864889A	2007-01-05 15:00:47.090000000	HR	36	
3	0EF5356002482715299901A	2012-10-22 11:02:35.010000000	NaN	36	
4	0F023589499656230C5E3E2	2013-09-14 18:38:39.097000000	NaN	36	

	LoanStatus	ClosedDate	BorrowerRate	ProsperScore	\
0	Completed	2009-08-14 00:00:00	0.1580	NaN	
1	Current	NaN	0.0920	7.0	
2	Completed	2009-12-17 00:00:00	0.2750	NaN	
3	Current	NaN	0.0974	9.0	
4	Current	NaN	0.2085	4.0	

	ListingCategory (numeric)	BorrowerState	Occupation	\
0	Not Available	CO	Other	
1	Home Improvement	CO	Professional	
2	Not Available	GA	Other	
3	Motorcycle	GA	Skilled Labor	
4	Home Improvement	MN	Executive	

	CreditScoreRangeLower	CreditScoreRangeUpper	DebtToIncomeRatio	\
0	640.0	659.0	0.17	
1	680.0	699.0	0.18	
2	480.0	499.0	0.06	
3	800.0	819.0	0.15	
4	680.0	699.0	0.26	

	IncomeRange	StatedMonthlyIncome	LoanOriginalAmount
0	\$25,000-49,999	3083.333333	9425
1	\$50,000-74,999	6125.000000	10000
2	Not displayed	2083.333333	3001

3	\$25,000-49,999	2875.000000	10000
4	\$100,000+	9583.333333	15000

2.0.3 Step 3. Remove null values from the 'CreditScoreRangeLower' column.

```
In [259]: df_score = df[df.CreditScoreRangeLower != 0]
```

Test

```
In [260]: df_score.describe()
```

```
Out[260]:
```

	Term	BorrowerRate	ProsperScore	CreditScoreRangeLower \
count	112933.000000	112933.000000	83982.000000	112342.000000
mean	40.805823	0.192907	5.954502	686.336544
std	10.426350	0.074928	2.373520	62.358444
min	12.000000	0.000000	1.000000	360.000000
25%	36.000000	0.134000	4.000000	660.000000
50%	36.000000	0.184000	6.000000	680.000000
75%	36.000000	0.250600	8.000000	720.000000
max	60.000000	0.497500	11.000000	880.000000

	CreditScoreRangeUpper	DebtToIncomeRatio	StatedMonthlyIncome \
count	112342.000000	104522.000000	1.129330e+05
mean	705.336544	0.276162	5.608873e+03
std	62.358444	0.553864	7.498445e+03
min	379.000000	0.000000	0.000000e+00
25%	679.000000	0.140000	3.204167e+03
50%	699.000000	0.220000	4.666667e+03
75%	739.000000	0.320000	6.833333e+03
max	899.000000	10.010000	1.750003e+06

	LoanOriginalAmount
count	112933.000000
mean	8321.863379
std	6237.048103
min	1000.000000
25%	4000.000000
50%	6350.000000
75%	12000.000000
max	35000.000000

2.0.4 Step 4. Change 'CreditScoreRangeLower', 'BorrowerRate' to integer.

```
In [261]: df.replace({'CreditScoreRangeLower': {None: 0}}, inplace=True )
df.CreditScoreRangeLower
```

```
Out[261]: 0      640.0
          1      680.0
          2      480.0
```

3	800.0
4	680.0
5	740.0
6	680.0
7	700.0
9	820.0
10	640.0
11	640.0
12	680.0
13	740.0
14	740.0
15	700.0
16	640.0
17	760.0
18	740.0
19	680.0
20	660.0
21	620.0
22	700.0
23	680.0
24	660.0
25	680.0
26	660.0
27	700.0
28	720.0
30	740.0
31	680.0
	...
113907	640.0
113908	700.0
113909	800.0
113910	640.0
113911	660.0
113912	800.0
113913	780.0
113914	520.0
113915	620.0
113916	660.0
113917	660.0
113918	740.0
113919	680.0
113920	740.0
113921	0.0
113922	640.0
113923	700.0
113924	640.0
113925	680.0
113926	540.0

```

113927    760.0
113928    740.0
113929    660.0
113930    680.0
113931    800.0
113932    700.0
113933    700.0
113934    700.0
113935    680.0
113936    680.0

```

Name: CreditScoreRangeLower, Length: 113066, dtype: float64

```
In [262]: df.ListingCreationDate
```

```

Out[262]: 0      2007-08-26 19:09:29.263000000
1      2014-02-27 08:28:07.900000000
2      2007-01-05 15:00:47.090000000
3      2012-10-22 11:02:35.010000000
4      2013-09-14 18:38:39.097000000
5      2013-12-14 08:26:37.093000000
6      2013-04-12 09:52:56.147000000
7      2013-05-05 06:49:27.493000000
9      2013-12-02 10:43:39.117000000
10     2012-05-10 07:04:01.577000000
11     2007-10-09 20:28:33.640000000
12     2013-12-15 20:01:10.757000000
13     2013-07-15 16:28:28.087000000
14     2013-04-19 11:17:41.700000000
15     2012-04-10 09:14:46.297000000
16     2013-07-16 12:42:48.680000000
17     2006-08-15 12:21:09.433000000
18     2013-02-20 03:48:37.470000000
19     2013-08-21 06:49:02.093000000
20     2013-11-22 11:35:02.987000000
21     2007-11-30 20:33:49.227000000
22     2013-01-30 09:36:13.783000000
23     2013-04-22 13:29:19.073000000
24     2013-12-03 11:34:46.127000000
25     2013-10-02 14:31:09.157000000
26     2013-02-12 18:07:31.690000000
27     2010-06-16 16:23:44.533000000
28     2013-11-02 16:01:28.050000000
30     2012-01-30 17:59:17.200000000
31     2012-09-21 13:37:43.210000000
...
113907 2013-11-09 06:55:05.690000000
113908 2010-11-29 10:40:09.730000000
113909 2012-09-08 10:34:38.837000000

```



```

113910    2014-01-24 19:56:49.960000000
113911    2013-11-15 11:42:47.540000000
113912    2013-12-16 16:36:00.990000000
113913    2010-04-25 15:13:27.963000000
113914    2006-08-09 14:34:40.010000000
113915    2008-07-29 05:22:29.390000000
113916    2012-11-08 20:07:36.600000000
113917    2013-11-23 04:52:50.057000000
113918    2008-06-19 12:02:53.300000000
113919    2013-05-07 18:49:59.750000000
113920    2013-06-11 05:49:40.247000000
113921    2005-11-09 20:44:28.847000000
113922    2008-08-08 16:58:54.760000000
113923    2008-09-10 08:26:30.537000000
113924    2012-10-20 19:15:52.670000000
113925    2013-04-25 13:54:45.017000000
113926    2006-08-01 10:31:31.143000000
113927    2008-04-30 21:25:19.670000000
113928    2011-06-06 19:02:44.443000000
113929    2013-07-06 17:40:01.657000000
113930    2013-07-08 10:24:49.700000000
113931    2014-01-16 20:13:08.040000000
113932    2013-04-14 05:55:02.663000000
113933    2011-11-03 20:42:55.333000000
113934    2013-12-13 05:49:12.703000000
113935    2011-11-14 13:18:26.597000000
113936    2014-01-15 09:27:37.657000000
Name: ListingCreationDate, Length: 113066, dtype: object

```

```

In [263]: df_int = df[df.ProspersScore != 0]
          df_int.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 113066 entries, 0 to 113936
Data columns (total 17 columns):
ListingKey                113066 non-null object
ListingCreationDate       113066 non-null object
CreditGrade              28953 non-null object
Term                     113066 non-null int64
LoanStatus               113066 non-null object
ClosedDate               55076 non-null object
BorrowerRate             113066 non-null float64
ProspersScore            83982 non-null float64
ListingCategory (numeric) 113066 non-null object
BorrowerState            107551 non-null object
Occupation               109537 non-null object
CreditScoreRangeLower    113066 non-null float64
CreditScoreRangeUpper    112475 non-null float64

```

```

DebtToIncomeRatio      104594 non-null float64
IncomeRange            113066 non-null object
StatedMonthlyIncome    113066 non-null float64
LoanOriginalAmount     113066 non-null int64
dtypes: float64(6), int64(2), object(9)
memory usage: 15.5+ MB

```

```
In [264]: df_int = df.astype({'CreditScoreRangeLower': 'int', 'BorrowerRate': 'int'})
```

```
In [265]: df_int.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 113066 entries, 0 to 113936
Data columns (total 17 columns):
ListingKey                113066 non-null object
ListingCreationDate       113066 non-null object
CreditGrade              28953 non-null object
Term                     113066 non-null int64
LoanStatus               113066 non-null object
ClosedDate               55076 non-null object
BorrowerRate             113066 non-null int64
ProsperScore              83982 non-null float64
ListingCategory (numeric) 113066 non-null object
BorrowerState            107551 non-null object
Occupation               109537 non-null object
CreditScoreRangeLower    113066 non-null int64
CreditScoreRangeUpper    112475 non-null float64
DebtToIncomeRatio        104594 non-null float64
IncomeRange              113066 non-null object
StatedMonthlyIncome      113066 non-null float64
LoanOriginalAmount       113066 non-null int64
dtypes: float64(4), int64(4), object(9)
memory usage: 15.5+ MB

```

2.0.5 Step 5. Remove values from 'Occupation' like "None", "Other", and null values

```
In [266]: df.dropna(subset=['Occupation'], inplace=True)
```

```
In [267]: df_occupation = df[(df['Occupation'] == 'Other')].index
df.drop(df_occupation, inplace=True)
```

Test

```
In [268]: df.Occupation
```

```

Out[268]: 1          Professional
          3          Skilled Labor

```

4	Executive
5	Professional
6	Sales - Retail
7	Laborer
9	Food Service
10	Fireman
11	Waiter/Waitress
12	Sales - Retail
13	Construction
14	Computer Programmer
16	Professional
17	Professional
18	Sales - Commission
19	Laborer
20	Retail Management
21	Professional
23	Skilled Labor
25	Engineer - Mechanical
26	Sales - Commission
27	Executive
28	Military Enlisted
32	Clerical
35	Retail Management
36	Professional
37	Teacher
43	Clergy
44	Professional
45	Executive
	...
113895	Engineer - Mechanical
113896	Food Service
113897	Doctor
113898	Sales - Retail
113899	Military Enlisted
113900	Accountant/CPA
113901	Executive
113903	Landscaping
113904	Nurse (LPN)
113906	Professional
113907	Sales - Retail
113908	Sales - Commission
113909	Clerical
113910	Executive
113912	Scientist
113913	Analyst
113916	Professional
113917	Clerical
113918	Social Worker

```

113920      Retail Management
113923      Clergy
113925      Homemaker
113927      Executive
113929      Accountant/CPA
113930      Professional
113931      Analyst
113932      Food Service Management
113933      Professional
113935      Food Service
113936      Professor
Name: Occupation, Length: 81115, dtype: object

```

2.0.6 Step 6: Create a random dataset

```

In [269]: np.random.seed(2018)
          sample = np.random.choice(df.shape[0], 300, replace = False)
          df_random2 = df.loc[sample]

```

/opt/conda/lib/python3.6/site-packages/ipykernel_launcher.py:3: FutureWarning:
Passing list-likes to .loc or [] with any missing label will raise
KeyError in the future, you can use .reindex() as an alternative.

See the documentation here:

<https://pandas.pydata.org/pandas-docs/stable/indexing.html#deprecate-loc-reindex-listlike>
This is separate from the ipykernel package so we can avoid doing imports until

```

In [270]: np.random.seed(2018)
          sample = np.random.choice(df.shape[0], 500, replace = False)
          df_random = df.loc[sample]

```

/opt/conda/lib/python3.6/site-packages/ipykernel_launcher.py:3: FutureWarning:
Passing list-likes to .loc or [] with any missing label will raise
KeyError in the future, you can use .reindex() as an alternative.

See the documentation here:

<https://pandas.pydata.org/pandas-docs/stable/indexing.html#deprecate-loc-reindex-listlike>
This is separate from the ipykernel package so we can avoid doing imports until

3 Analyze

3.1 Univariate Exploration

3.1.1 Average interest rate of most Borrowers

```

In [271]: df.describe()

```

```

Out[271]:
      Term  BorrowerRate  ProsperScore  CreditScoreRangeLower \
count  81115.000000  81115.000000  61586.000000  81115.000000
mean    40.929594    0.191218    6.037720    687.528324
std     10.568458    0.075064    2.388946    64.839848
min     12.000000    0.000000    1.000000    0.000000
25%     36.000000    0.131400    4.000000    660.000000
50%     36.000000    0.181400    6.000000    680.000000
75%     36.000000    0.249900    8.000000    720.000000
max     60.000000    0.360000   11.000000   880.000000

      CreditScoreRangeUpper  DebtToIncomeRatio  StatedMonthlyIncome \
count      81114.000000      75977.000000      8.111500e+04
mean        706.536800        0.262449      5.970960e+03
std          64.795293        0.461642      7.858231e+03
min          19.000000        0.000000      0.000000e+00
25%         679.000000        0.140000      3.416667e+03
50%         699.000000        0.220000      5.000000e+03
75%         739.000000        0.310000      7.166667e+03
max          899.000000       10.010000     1.750003e+06

      LoanOriginalAmount
count      81115.000000
mean       8612.301633
std        6377.056491
min        1000.000000
25%        4000.000000
50%        7000.000000
75%       12500.000000
max       35000.000000

```

This chart shows that the average interest rate in this data set is just about 19 percent, whereas, some people have 0 percent interest rates and some have up to 50 percent interest rates but those are pretty rare. It would be interesting to see the correlation between the income range and interest rate

3.1.2 Correlation between Stated Monthly Income and the Borrower Rate

3.1.3 The correlation between interest rate and credit score

```
In [272]: df.describe()
```

```

Out[272]:
      Term  BorrowerRate  ProsperScore  CreditScoreRangeLower \
count  81115.000000  81115.000000  61586.000000  81115.000000
mean    40.929594    0.191218    6.037720    687.528324
std     10.568458    0.075064    2.388946    64.839848
min     12.000000    0.000000    1.000000    0.000000
25%     36.000000    0.131400    4.000000    660.000000
50%     36.000000    0.181400    6.000000    680.000000
75%     36.000000    0.249900    8.000000    720.000000

```

max	60.000000	0.360000	11.000000	880.000000
-----	-----------	----------	-----------	------------

	CreditScoreRangeUpper	DebtToIncomeRatio	StatedMonthlyIncome \
count	81114.000000	75977.000000	8.111500e+04
mean	706.536800	0.262449	5.970960e+03
std	64.795293	0.461642	7.858231e+03
min	19.000000	0.000000	0.000000e+00
25%	679.000000	0.140000	3.416667e+03
50%	699.000000	0.220000	5.000000e+03
75%	739.000000	0.310000	7.166667e+03
max	899.000000	10.010000	1.750003e+06


	LoanOriginalAmount
count	81115.000000
mean	8612.301633
std	6377.056491
min	1000.000000
25%	4000.000000
50%	7000.000000
75%	12500.000000
max	35000.000000

When the individual pulled their credit score, they get an upper and lower credit score. On average, the lower credit score was 686 and the lower is 705. I will use this to explore how the lower score correlates with the interest rate.

3.1.4 Most common loan types

```
In [273]: base_color = sb.color_palette()[0]
ListingCategory = df['ListingCategory (numeric)'].value_counts().index
sb.countplot(data = df, x = 'ListingCategory (numeric)', color = base_color)
plt.xticks(rotation = 90)
```

```
Out[273]: (array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
17, 18, 19, 20]), <a list of 21 Text xticklabel objects>)
```



By far, the most common loan type is the debt consolidation loan. It would be interesting to see the correlation between loan categorie and loan amount.

3.1.5 What is the highest, lowest and the average risk scores

In [274]: `df.describe()`

```
Out[274]:
```

	Term	BorrowerRate	ProsperScore	CreditScoreRangeLower \
count	81115.000000	81115.000000	61586.000000	81115.000000
mean	40.929594	0.191218	6.037720	687.528324
std	10.568458	0.075064	2.388946	64.839848
min	12.000000	0.000000	1.000000	0.000000
25%	36.000000	0.131400	4.000000	660.000000
50%	36.000000	0.181400	6.000000	680.000000
75%	36.000000	0.249900	8.000000	720.000000
max	60.000000	0.360000	11.000000	880.000000

	CreditScoreRangeUpper	DebtToIncomeRatio	StatedMonthlyIncome \
--	-----------------------	-------------------	-----------------------

count	81114.000000	75977.000000	8.111500e+04
mean	706.536800	0.262449	5.970960e+03
std	64.795293	0.461642	7.858231e+03
min	19.000000	0.000000	0.000000e+00
25%	679.000000	0.140000	3.416667e+03
50%	699.000000	0.220000	5.000000e+03
75%	739.000000	0.310000	7.166667e+03
max	899.000000	10.010000	1.750003e+06


	LoanOriginalAmount
count	81115.000000
mean	8612.301633
std	6377.056491
min	1000.000000
25%	4000.000000
50%	7000.000000
75%	12500.000000
max	35000.000000

The highest risk score is 11 and the lowest score is 1 (out of 10). The average risk score is about 6. It would be interesting to see if there is a correlation between risk score and monthly income.

3.2 Which state has the highest borrowing rate

```
In [275]: plt.figure(figsize=(20,12))
          base_color = sb.color_palette()[0]
          ListingCategory = df['BorrowerState'].value_counts().index
          sb.countplot(data = df, x = 'BorrowerState', color = base_color)
          plt.xticks(rotation = 90)

Out[275]: (array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
                  17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33,
                  34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50]),
          <a list of 51 Text xticklabel objects>)
```





The most popular borrower state is California which is not surprising because they have a high population. It is not surprising that the other states are higher. Many of the top borrowing states have high population. It would be worth looking at each state versus population.

3.2.1 Difference between occupation and how much they borrowed

```
In [276]: plt.figure(figsize=(20,12))
          base_color = sb.color_palette()[0]
          ListingCategory = df['Occupation'].value_counts().index
          sb.countplot(data = df, x = 'Occupation', color = base_color)
          plt.xticks(rotation = 90)

Out[276]: (array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
                  17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33,
                  34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50,
                  51, 52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64, 65]),
          <a list of 66 Text xticklabel objects>)
```



The most common occupation for borrowers is "Professional". Since this is so vague, I would like to remove this term and look more into the correlation between occupation and loan amount.

```
In [277]: df_occupation = df[(df['Occupation'] == 'Professional')].index
df.drop(df_occupation, inplace=True)
df['Occupation'].value_counts()
```


```
Out[277]: Computer Programmer      4442
Executive      4280
Teacher      3729
Administrative Assistant      3677
Analyst      3578
Sales - Commission      3414
Accountant/CPA      3209
Clerical      3135
Sales - Retail      2780
Skilled Labor      2723
Retail Management      2579
Nurse (RN)      2475
Construction      1777
Truck Driver      1666
```

Laborer	1578
Police Officer/Correction Officer	1561
Civil Service	1445
Engineer - Mechanical	1397
Military Enlisted	1265
Food Service Management	1227
Engineer - Electrical	1118
Food Service	1115
Medical Technician	1105
Attorney	1033
Tradesman - Mechanic	937
Social Worker	733
Postal Service	622
Professor	556
Realtor	535
Doctor	492
	...
Scientist	367
Military Officer	341
Bus Driver	314
Principal	307
Teacher's Aide	275
Pharmacist	255
Student - College Graduate Student	245
Landscaping	234
Engineer - Chemical	222
Investor	213
Architect	212
Pilot - Private/Commercial	198
Clergy	195
Student - College Senior	187
Car Dealer	177
Chemist	144
Psychologist	143
Biologist	125
Religious	124
Flight Attendant	122
Tradesman - Carpenter	119
Homemaker	119
Student - College Junior	111
Tradesman - Plumber	102
Student - College Sophomore	69
Dentist	67
Student - College Freshman	41
Student - Community College	28
Judge	22
Student - Technical School	16

Name: Occupation, Length: 65, dtype: int64

```
In [278]: plt.figure(figsize=(20,12))
          base_color = sb.color_palette()[0]
          ListingCategory = df['Occupation'].value_counts().index
          sb.countplot(data = df, x = 'Occupation', color = base_color)
          plt.xticks(rotation = 90)

Out[278]: (array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
                    17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33,
                    34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50,
                    51, 52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64]),
          <a list of 65 Text xticklabel objects>)
```



In this graph, computer programmers, executives, teachers and administrative assistants borrow the most frequently.

3.2.2 Discuss the distribution(s) of your variable(s) of interest. Were there any unusual points? Did you need to perform any transformations?

There were a few changes that needed to be made throughout the exploring process. Personally, I find it easier to complete these cleaning tasks (besides the major ones) as I go. Cleaning steps 3 through 5 were completed after the analysis began.

3.2.3 Of the features you investigated, were there any unusual distributions? Did you perform any operations on the data to tidy, adjust, or change the form of the data? If so, why did you do this?

There were a few instances where I believed that the data would've been more defined if values were removed. In the 'difference between occupation and how much they borrowed' portion of the assessment, I found that the occupation "Professional" was wildly used. Beings this isn't specific, I believe it would be best to remove so you can clearly see which occupations were the most popular.


3.3 Bivariate Exploration

In this section, I will investigate relationships between pairs of variables in your data.

```
In [279]: plt.figure(figsize = [20, 10])
```

```
bins_x = np.arange(0, 14000.5+1, 500)
bins_y = np.arange(0, 0.4+.1, .025)
```

```
plt.hist2d(data = df, x = 'StatedMonthlyIncome', y = 'BorrowerRate', bins = [bins_x, b
plt.xlabel('Monthly Income')
plt.ylabel('Interest Rate')
plt.colorbar();
```



The graph shows that the majority of borrowers make between 4,000 and 4,500 dollars monthly and have an interest rate between 15 and 20 percent. The majority of borrowers make between 2,500 and 5,000 dollars monthly.


3.3.1 Correlation between monthly income and prosper score

```
In [280]: plt.figure(figsize = [20, 10])

bins_x = np.arange(0, 14000.5+1, 500)
bins_y = np.arange(0, 11.0+.1, 1)

plt.hist2d(data = df, x = 'StatedMonthlyIncome', y = 'ProsperScore', bins = [bins_x, b
plt.xlabel('Monthly Income')
plt.ylabel('ProsperScore')
plt.colorbar();

/opt/conda/lib/python3.6/site-packages/numpy/lib/function_base.py:968: RuntimeWarning: invalid v
not_smaller_than_edge = (sample[:, i] >= edges[i][-1])
```



There is not a huge correlation between monthly income and prosper score. There is quite a bit of yellow between 3,000 dollars monthly income and 5,500 dollars monthly income and the prosper score ranges from 4 to 9.

This graph shows that the majority of people have credit scores between 700 and 725. Those credit scores have interest rates between 10 and 15 percent.

3.3.2 Correlation between occupation and how much they borrowed


```
In [281]: plt.figure(figsize=(20,12))

base_color = sb.color_palette()[0]
```

```
ax = sb.barplot(data = df, x = 'Occupation', y = 'StatedMonthlyIncome', color = base_color)
ax.legend(loc = 20, ncol = 3, framealpha = 1, title = 'cat_var2')
```

```
plt.xticks(rotation = 90)
```

```
Out[281]: (array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
        17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33,
        34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50,
        51, 52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64]),
  <a list of 65 Text xticklabel objects>)
```



According to the graph, doctors, attorneys, dentists, judges and executives borrow the most amount of money.


```
In [282]: plt.figure(figsize=(20,12))
```

```
base_color = sb.color_palette()[0]
```

```
ax = sb.barplot(data = df, x = 'Occupation', y = 'LoanOriginalAmount', color = base_color)
ax.legend(loc = 20, ncol = 3, framealpha = 1, title = 'cat_var2')
```

```
plt.xticks(rotation = 90)
```

```
Out[282]: (array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
        17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33,
        34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50,
        51, 52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64]),
        <a list of 65 Text xticklabel objects>)
```




Judges, pharmacists, doctors and executives make the most amount of money. Whereas, judges, pharmacists, doctors, executives, and pilots borrow the most amount of money. It makes sense, if you have a lot of money, you have more money to borrow.

3.3.3 The correlation between interest rate and credit score

```
In [283]: plt.figure(figsize = [15, 5])
```

```
bins_x = np.arange(500, 800.5+1, 25)
bins_y = np.arange(0, 0.5+.1, .025)
```

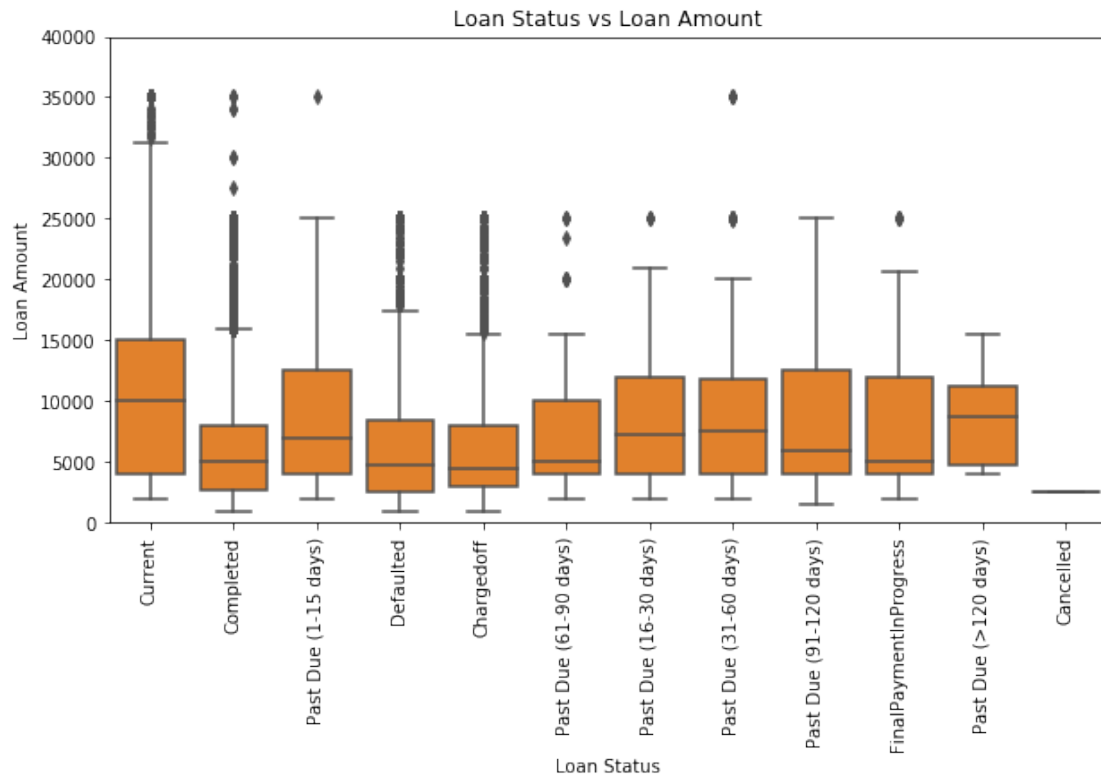
```
plt.hist2d(data = df, x = 'CreditScoreRangeLower', y = 'BorrowerRate', bins = [bins_x,
plt.xlabel('Credit Score')
plt.ylabel('Interest Rate')
plt.colorbar();
```

This graph shows that the majority of borrowers had a credit score of between 700 and 725 and their interest rates fell around 15%.

3.3.4 Loan status versus loan amount

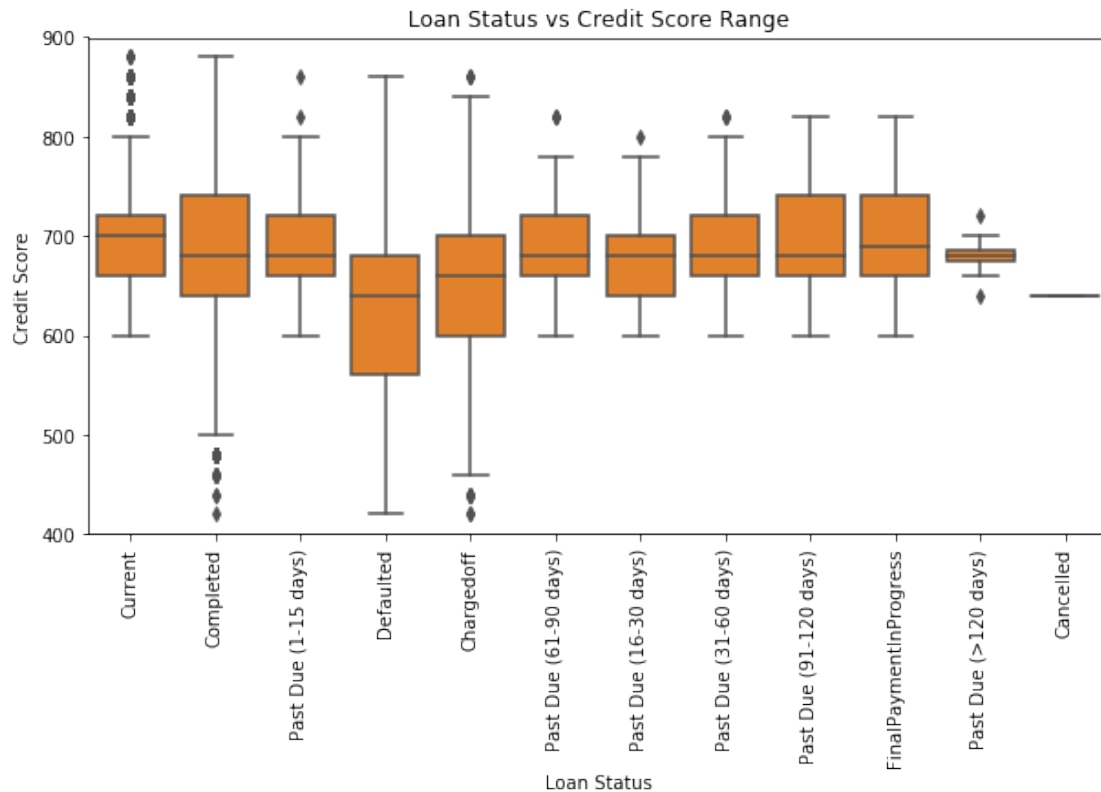
```
In [284]: plt.figure(figsize = [10, 5])
          base_color = sb.color_palette()[1]
          sb.boxplot(data = df, x = 'LoanStatus', y = 'LoanOriginalAmount', color = base_color)
          plt.ylim([0, 40000])
          plt.title('Loan Status vs Loan Amount')
          plt.xlabel('Loan Status')
          plt.ylabel('Loan Amount')
          plt.xticks(rotation = 90)
          plt.show()
```



Loans that are in the 10,000 dollar range are more likely to be current, whereas, seriously past due loans (91-120 days) are about 5,000. Many loans that have defaulted are lower loan amounts (on average, about 5,000 dollars).

3.3.5 Loan Status versus credit score

```
In [285]: plt.figure(figsize = [10, 5])
          base_color = sb.color_palette()[1]
          sb.boxplot(data = df, x = 'LoanStatus', y = 'CreditScoreRangeLower', color = base_color)
          plt.ylim([400, 900])
          plt.title('Loan Status vs Credit Score Range')
          plt.xlabel('Loan Status')
          plt.ylabel('Credit Score')
          plt.xticks(rotation = 90)
          plt.show()
```




This graph shows that people with 700 and higher credit scores are much more likely to pay for their loan. This graph also shows that if people have a 650 credit score or lower, they are more likely to default on their loan.

3.3.6 Loan category and the amount of the loan

```
In [286]: plt.figure(figsize=(20,12))
           base_color = sb.color_palette()[0]
           sb.boxplot(data = df, x = 'ListingCategory (numeric)', y = 'LoanOriginalAmount', color=
           plt.xticks(rotation = 90)
```

```
Out[286]: (array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
                  17, 18, 19, 20]), <a list of 21 Text xticklabel objects>)
```



This graph shows that people that are borrowing for baby & adoption tend to borrow much more than other categories. Debt consolidation is the most popular loan and they are second highest borrower.

3.3.7 Talk about some of the relationships you observed in this part of the investigation. How did the feature(s) of interest vary with other features in the dataset?

There is a strong correlation between interest rate and credit score is high. Most of borrowers fall into a particular interest rate and credit score. Secondly, there is an interesting correlation between occupation and many different factors including how much their monthly income is, the loan amount, etc. The relationship between income, occupation and loan amount was not shocking; the more money someone made, the more they spent. The correlation between interest rate and credit score was the most interesting relationship that I found.

One of the graph shows that the majority of borrowers make between 4,000 and 4,500 dollars monthly and have an interest rate between 15 and 20 percent. The majority of borrowers make between 2,500 and 5,000 dollars monthly.

Another graph shows that people with 700 and higher credit scores are much more likely to pay for their loan. This graph also shows that if people have a 650 credit score or lower, they are more likely to default on their loan.

By far, the most common loan type is the debt consolidation loan.

3.3.8 Did you observe any interesting relationships between the other features (not the main feature(s) of interest)?

I found some interesting data while exploring how many people work at a certain occupation got loans and also the relationship between occupation and loan amount.


Judges, pharmacists, doctors and executives borrow the most amount of money. Whereas, computer programmers, executives, teachers and administrative assistants borrow more frequently

3.4 Multivariate Exploration

3.4.1 Borrowing Rate versus Monthly Income versus Credit Grade

```
In [287]: g = sb.FacetGrid(data = df_random, hue = "CreditGrade", hue_order = ['AA', 'A', 'B', 'C', 'D', 'E', 'HR'],
g.map(sb.regplot, "BorrowerRate", "StatedMonthlyIncome", fit_reg = False);

new_title = 'Credit Grade'
g.add_legend(title = 'Credit Grade')
plt.ylabel('Interest Rate')
plt.xlabel('Monthly Income')
plt.title("Interest Rate vs Monthly Income and Credit Grade");
```



According to the graph, borrowers that had an 'AA' credit score, also had low interest rates and their income ranged anywhere from about 16,000 dollars a month to 0 dollars per month. As you can see, the interest rate and the credit score are directly related and you can see that as the credit interest rate gets higher, the monthly income comes down.

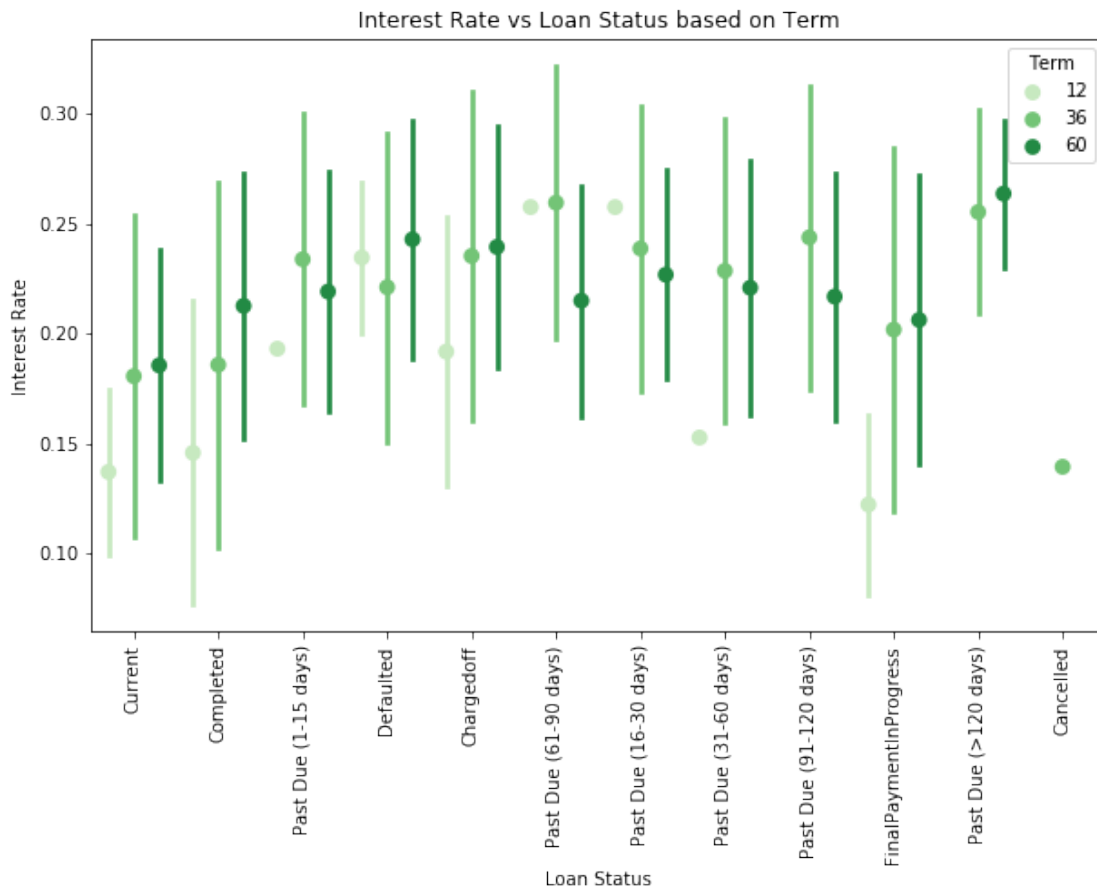
3.5 Interest Rate vs Loan Status based on Term

```
In [288]: fig = plt.figure(figsize = [10,6])
ax = sb.pointplot(data = df, x = 'LoanStatus', y = 'BorrowerRate', hue = 'Term',
```

```

        palette = 'Greens', linestyle = '', dodge = .6, ci='sd')
plt.title('Interest Rate vs Loan Status based on Term')
plt.ylabel('Interest Rate')
plt.xlabel('Loan Status')
plt.xticks(rotation = 90)
ax.set_yticklabels([],minor = True);

```



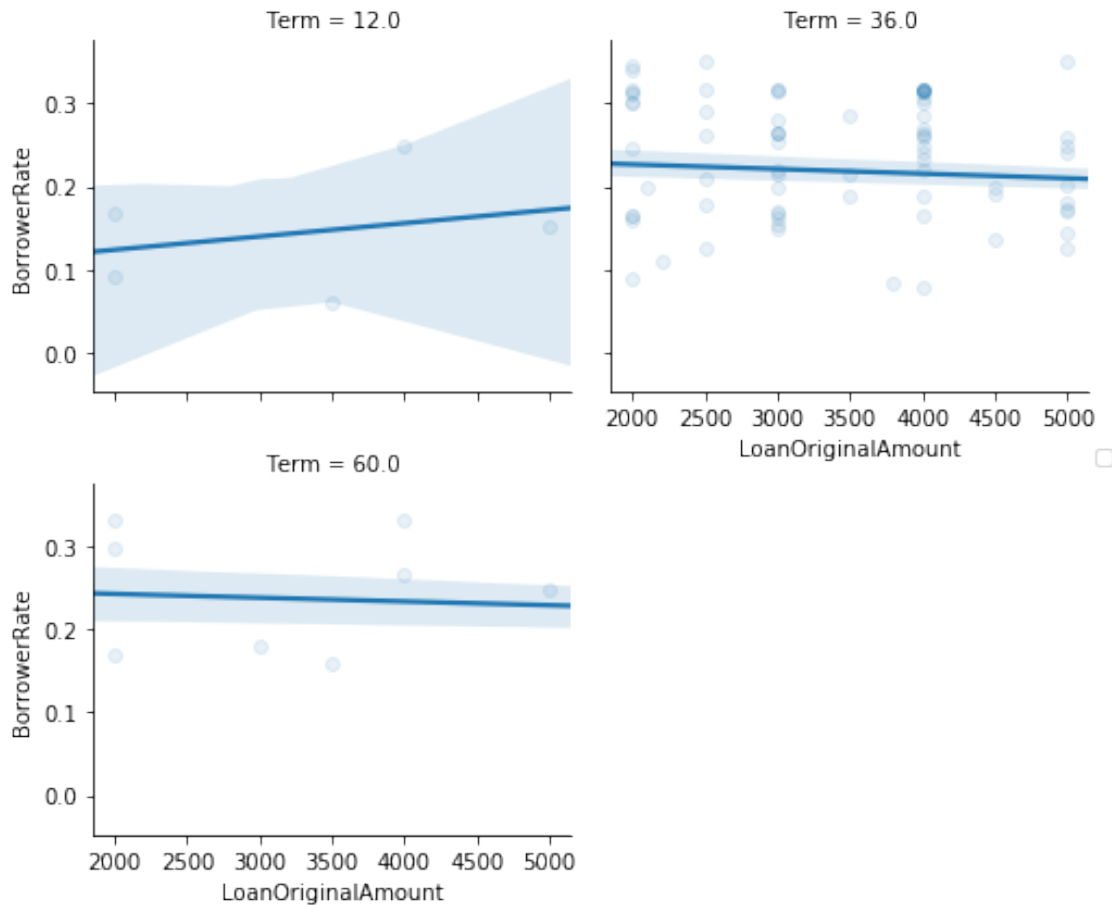
This graph shows that current loans that have a shorter term has less of an interest rate. This is the same for completed. Surprisingly, the past due loan status' typically have higher interest rates. This would make sense - based on my analysis, people that have bad credit scores are at high risk (not paying back their loan). If someone has a high interest rate, they most likely have a bad credit score and thus higher interest rates.

3.5.1 Loan Amount versus Interest Rate and Loan Term

```

In [290]: g=sb.FacetGrid(data=df_random2, aspect=1.2, col='Term', col_wrap=2)
          g.map(sb.regplot, 'LoanOriginalAmount', 'BorrowerRate', x_jitter=0.04, scatter_kws={'a
          g.add_legend();

```




In a 12 month term, it appears that there is a slight increase in interest rate the more the loan is. In a 36 month term, there is a slight (barely noticable) decline in interest rate the more money that was borrowed.

3.6 Loan Amount versus Monthly Income and Term

```
In [291]: g = sb.FacetGrid(data = df_random, hue = 'Term', size = 10,
                        palette = 'mako_r')
g.map(plt.scatter, 'StatedMonthlyIncome', 'LoanOriginalAmount')
plt.ylabel('Monthly Income')
plt.xlabel('Loan Amount')
g.add_legend()
```

```
Out[291]: <seaborn.axisgrid.FacetGrid at 0x7f22f01f2c88>
```



This graph shows that the majority of people making between 1,000 and 10,000 a month that also borrowed between 0 to 8,000 dollars typically have a 36 month loan. It is interesting looking at the 60 month loan because they are typically always either (about) make around 4,000, 1,000 and 15,000 dollars a month. It appears to be lines in the graph right where those income ranges are. Many of the 60 month loans are less than 10,000 dollars. Most of the loans are 12 months.

3.6.1 Talk about some of the relationships you observed in this part of the investigation. Were there features that strengthened each other in terms of looking at your feature(s) of interest?

Current loans that have a shorter term has less of an interest rate. This is the same for completed. Surprisingly, the past due loan status' typically have higher interest rates. This would make sense - based on my analysis, people that have bad credit scores are at high risk (not paying back their loan). If someone has a high interest rate, they most likely have a bad credit score and thus higher interest rates.

Borrowers that had an 'AA' credit score, also had low interest rates and their income ranged anywhere from about 16,000 dollars a month to 0 dollars per month. As you can see, the interest rate and the credit score are directly related and you can see that as the credit interest rate gets higher, the monthly income comes down.

In a 12 month term, it appears that there is a slight increase in interest rate the more the loan is. In a 36 month term, there is a slight (barely noticable) decline in interest rate the more money that was borrowed.

The majority of people making between 1,000 and 10,000 a month that also borrowed between 0 to 8,000 dollars typically have a 36 month loan. It is interesting looking at the 60 month loan because they are typically always either (about) make around 4,000, 1,000 and 15,000 dollars a month. It appears to be lines in the graph right where those income ranges are. Many of the 60 month loans are less than 10,000 dollars. Most of the loans are 12 months.

3.6.2 Were there any interesting or surprising interactions between features?

The most surprising interaction in this dataset is the correlation between borrowing rate versus monthly income versus credit grade because of how clear the graph shows that monthly income, interest rate and the credit grade are directly linked and it shows perfectly on the map.

In []: